

The QUIC transport protocol: design and Internet-Scale deployment

David Calavera
CTO, Netlify



- **Motivation**
- **Design and implementation**
- **Internet-Scale deployment**
- **Performance**
- **Lessons learned**
- **Conclusion**

- Motivation
- Design and implementation
- **Testing in Production FTW**
- Performance
- Lessons learned
- Conclusion

Motivation



How does the internet work?

Welcome to my favourite job interview question!

TCP+TLS Handshake Delay

Sooooo chatty

David
wants to see
pictures of
vegan food

David's
Computer

The
Internet



David's
Computer

The
Internet




Or, I can show
David pictures
of vegan food

David's
Computer

The
Internet

David doesn't
want anyone to
see the pictures
of vegan food
he's looking at



David's
Computer


The
Internet



OK, I can make sure
nobody sees the pictures
of vegan food that
David is looking at

David's
Computer


The
Internet



Please, use
this key to
make sure nobody
sees the pictures
of vegan food that
David is looking at

David's
Computer

The
Internet



OK, I'll use that
key to make sure that
nobody sees the pictures
of vegan food that
David is looking at

**All Internet traffic
is going to be
encrypted sooner
rather than later**





Figure 3: Increase in secure web traffic to Google's front-end servers.

Percentage of Web Pages Loaded by Firefox Using HTTPS


(14-day moving average, source: [Firefox Telemetry](#))



Head-of-line blocking

Sooooo bossy

David's
computer



The
Internet

TCP Connection

**TCP is virtually
impossible to
change**

Design and implementation

An abstract geometric pattern in the bottom right corner of the slide. It consists of several teal-colored triangles of varying sizes, some of which are outlined with dark lines, creating a complex, crystalline structure against the dark blue background.




Figure 1: QUIC in the traditional HTTPS stack.

Connection establishment

0 Round-Trip Time

* At best

David wants to see pictures of vegan food. He doesn't want anyone to see the pictures of vegan food he's looking at. He's been here before, use this key to ensure that nobody sees the pictures of vegan food that he's looking at.

David's
Computer

The
Internet



David's
Computer

The
Internet



Cool.

3.1 Connection Establishment

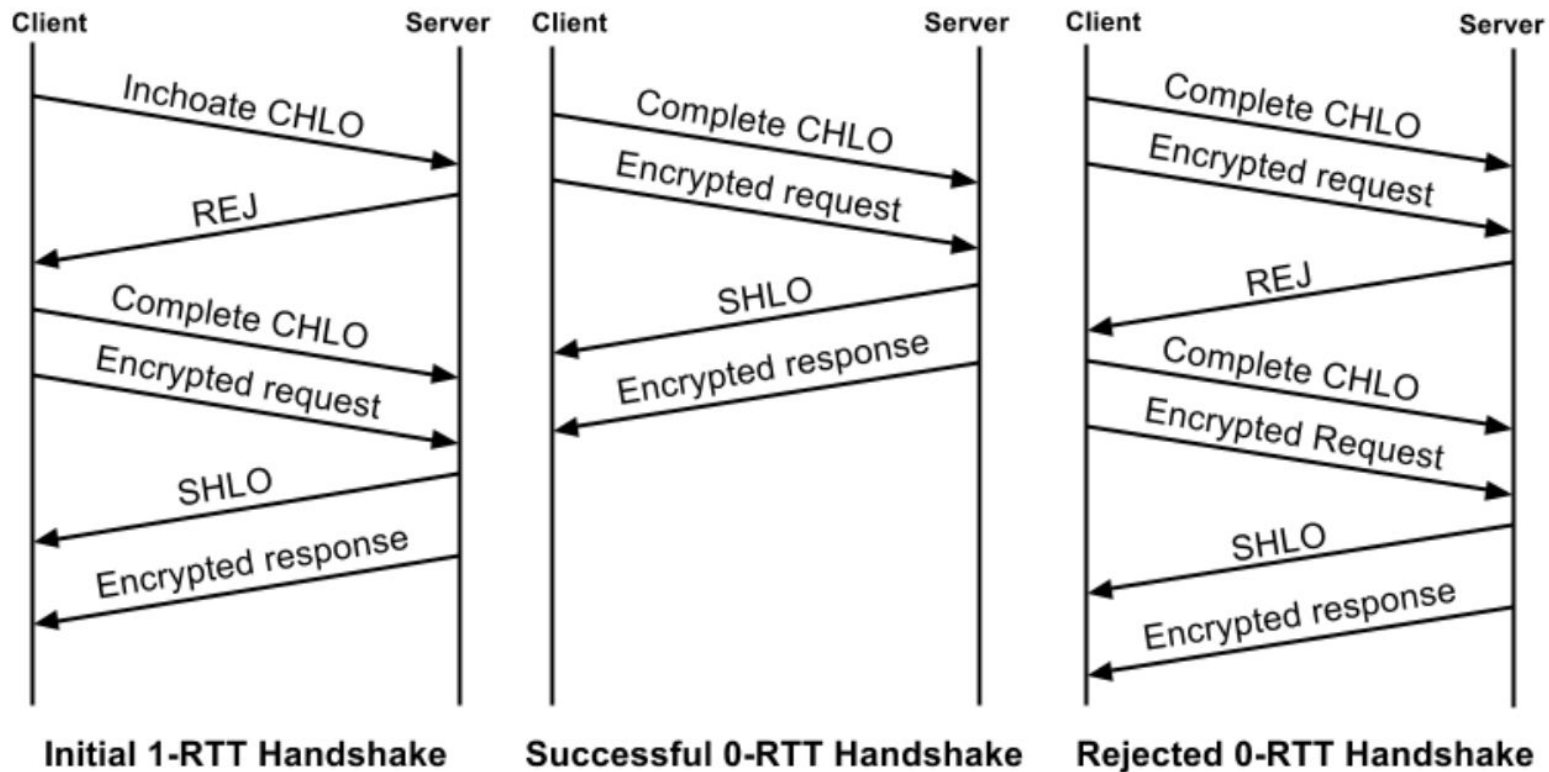


Figure 4: Timeline of QUIC's initial 1-RTT handshake, a subsequent successful 0-RTT handshake, and a failed 0-RTT handshake.

Stream multiplexing

**"A lost UDP packet
only impacts
streams whose
data it carries"**

Authentication & Encryption

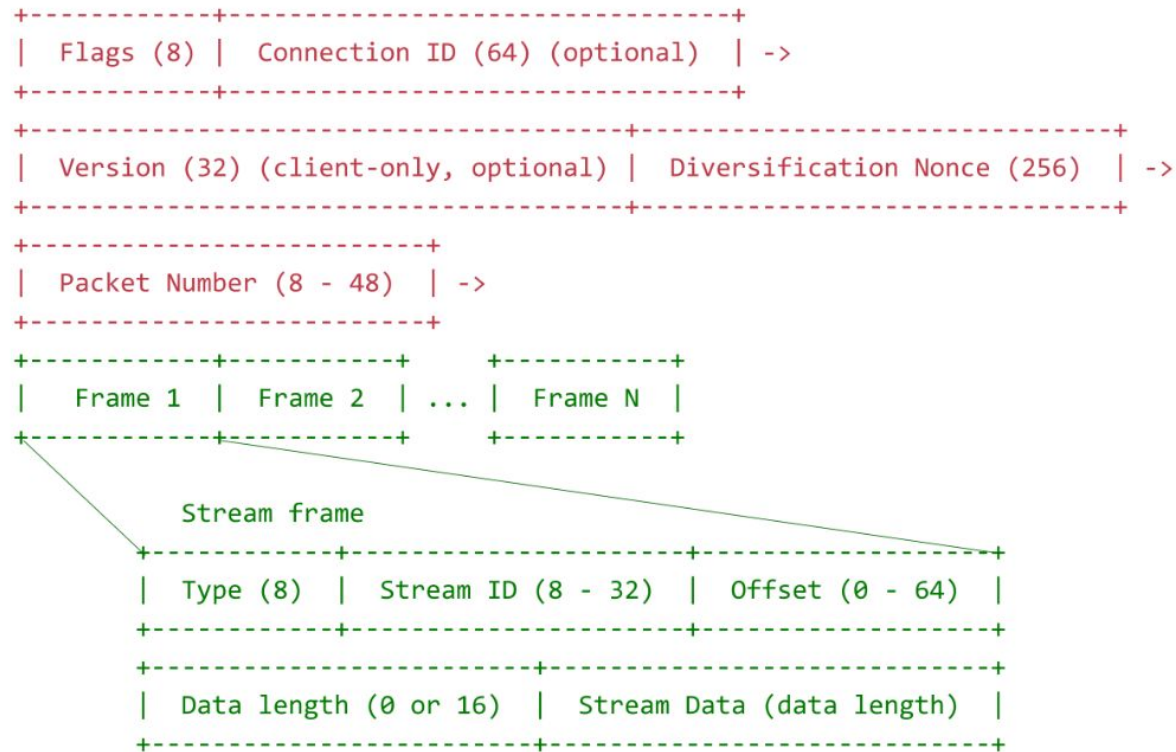


Figure 5: Structure of a QUIC packet, as of version 35 of Google's QUIC implementation. Red is the authenticated but unencrypted public header, green indicates the encrypted body. This packet structure is evolving as QUIC gets standardized at the IETF [2].

Loss recovery

**"Each packet
carries a new
monotonic
number"**

**"Acknowledgments
encode the delay
since receiving the
packet"**

Congestion control

Roll your own

NAT rebinding & Connection migration

**"QUIC connections
are identified
by a 64-bit
connection ID"**

Internet-Scale deployment



**"QUIC support was
added to Chrome in
June 2013"**

**"Allowed to new
features to be A/B
tested"**




Figure 6: Search Latency reduction for users in the QUIC experiment over an 18-month period. Numbered events are described in Section 5.2.

Performance



Transport and application metrics




Figure 7: Comparison of handshake latency for QUIC_g and TCP_g versus the minimum RTT of the connection. Solid lines indicate the mean handshake latency for all connections, including 0-RTT connections.

		% latency reduction by percentile						
		Lower latency				Higher latency		
	Mean	1%	5%	10%	50%	90%	95%	99%
Search								
Desktop	8.0	0.4	1.3	1.4	1.5	5.8	10.3	16.7
Mobile	3.6	-0.6	-0.3	0.3	0.5	4.5	8.8	14.3
Video								
Desktop	8.0	1.2	3.1	3.3	4.6	8.4	9.0	10.6
Mobile	5.3	0.0	0.6	0.5	1.2	4.4	5.8	7.5

Table 1: Percent reduction in global Search and Video Latency for users in QUIC_g, at the mean and at specific percentiles. A 16.7% reduction at the 99th percentile indicates that the 99th percentile latency for QUIC_g is 16.7% lower than the 99th percentile latency for TCP_g.

		% rebuffer rate reduction by percentile				
		Fewer rebufferers		More rebufferers		
	Mean	< 93%	93%	94 %	95%	99%
Desktop	18.0	*	100.0	70.4	60.0	18.5
Mobile	15.3	*	*	100.0	52.7	8.7

Table 2: Percent reduction in global Video Rebuffer Rate for users in QUIC_g at the mean and at specific percentiles. An 18.5% reduction at the 99th percentile indicates that the 99th percentile rebuffer rate for QUIC_g is 18.5% lower than the 99th percentile rate for TCP_g. An * indicates that neither QUIC_g nor TCP_g have rebufferers at that percentile.

Server CPU utilization

**"QUIC's server
CPU utilization
was about 3.5
times higher"**

Lessons learned



Packet size considerations





Figure 12: Unreachability with various UDP payload sizes. Data collected over 28 days in January 2014.

UDP Blockage and Throttling

**"QUIC is
successfully for
95.3% of client
connections"**

Forward Error Correction



"The benefits are limited to under 30% of loss episodes"

User-Space Development

**"Led us to uncover
a decade old
Cubic bug"**

Experiences with Middleboxes

RTFS*

* Read The F***** Specification

Conclusion



How Secure and Quick is QUIC

R. Lychev, S. Jero, A. Boldyreva, and C. Nita-Rotaru. 2015.

**"Does NOT satisfy
the traditional
notion of forward
secrecy"**

Standard WIP

<https://quicwg.github.io>

**Not widely
implemented outside
Google yet**

Chromium, Apache Traffic Server, Caddy

Thank you for listening!

David Calavera
@calavera

